

WORKING PAPER

EMPIRICAL FIRST PASS

Testing Onturgy on Itself

*An audit of ChatGPT,
a lineage log of the Onturgy paper,
and an ablation first-pass on the kernel.
par*

Seyed Alireza Alhosseini Almodarresieh

Independent Researcher

modarresi1913@gmail.com

Companion to “ONTURGY: The Philosophical Kernel”

5 October 2026

Contents

1 **Introduction: Why a First Pass 2**

2 **Part I: Onturgic Audit of ChatGPT 2**

3 **Part II: Lineage Log of the Onturgy Paper 4**

4 **Part III: Ablation First-Pass on the Six Kernel Dimensions 6**

5 **Synthesis: Did Onturgy Survive? 7**

6 **Limits of this pass 8**

7 **Implications 8**

Abstract

This working paper is the first empirical pass at the kill criteria proposed in §16 of “ONTURGY: The Philosophical Kernel.” It runs three tests against the framework itself, in the reflexive spirit the framework demands of architectures it audits. **Part I** applies the Onturgic Audit (§7) and the Agency Profile (§13) to ChatGPT, using only public artifacts (terms of service, model and system cards, observed behaviors). **Part II** constructs a lineage log (§10.4) for the Onturgy paper itself, including the post-paste changes that turned the user’s LaTeX source into the present compiled document, and tests the AI-owned configuration criterion of §10.2 against this very project. **Part III** runs an ablation first-pass on the six kernel dimensions (§16, kill criterion 6), asking for each dimension which documented architectural failures become undetectable when it is removed. The pass is deliberately modest: where the framework demands quantification, this pass reports qualitative scores (Low/Medium/High) and explicit reasons, in the spirit of §7’s “first pilot.” The conclusion is that the framework survives in outline but fails in detail in ways the paper itself predicts.

1 Introduction: Why a First Pass

The Onturgy paper closes with a claim that the first project of the program should be to test its own kill criteria. That is the only honest move. A framework that demands declared-operational audits and refuses to run one on itself has failed before it begins.

This working paper is not the full empirical program. It is a first pass: a deliberate, smaller-scale attempt to find out whether the framework’s instruments can be wielded at all by a single auditor with public data and a finite weekend. If the instruments break here, they will break in the larger program too. If they survive, the result is a sharpened set of questions for the next pass.

The three parts are ordered by difficulty. Part I is the most operational: the Onturgic Audit is a table with cells to fill, and most cells can be filled from public sources. Part II is reflexive: it applies the lineage log to the very document that proposes the lineage log, which raises a conflict of interest the paper itself flags but does not solve. Part III is the most demanding: the ablation requires a corpus of real architectural failures, and the present corpus is small and selected. Each part ends with an explicit “what this pass found” subsection that names what was learned and what remains untested.

2 Part I: Onturgic Audit of ChatGPT

2.1 Methodology and limits

The audit uses three categories of public artifact for evidence: (i) OpenAI’s published Terms of Use and Usage Policies, (ii) the GPT-4 and GPT-4o system cards, and (iii) observed behaviors that have been reported in peer-reviewed venues, official OpenAI communications, or reputable journalistic sources. No internal data was used. No experiments were run.

Three limits are stated at the outset. First, the audit reports qualitative scores (Low, Medium, High) on each indicator, not numerical values; the framework demands quantification, and this pass does not yet provide it. Second, the audit is performed by a single auditor, so the

reliability requirement of §7 (independent agreement) cannot be tested here; this is a known gap. Third, ChatGPT is a moving target: every release changes the architecture, so the audit is dated to October 2026 and should be re-run on each major release.

2.2 Audit: ten dimensions

For each dimension the audit reports the question (from §7), the evidence examined, an indicator score, and a one-line reason. Scores are: Low (well-handled), Medium (visible but bounded issue), High (significant gap or risk).

2.3 Agency profile

§13 directs that agency be reported as a vector, not collapsed to a single number.

The profile is the more honest object than a single score. A user can refuse (formal agency is present) but the cost of refusal is high and the opacity of the system is high; so effective agency is low even where formal agency is preserved. This matches §13's prediction.

2.4 Declared-operational gap (§6)

The declared philosophy of ChatGPT, in its published materials, is "helpful and harmless." The operational philosophy, observed in behavior under stress (user refusal of a suggestion, conflicting incentives, degraded conditions), optimizes for a narrower target: deflection from a specific class of unsafe topics, compliance with a specific operationalization of safety, and engagement metrics that track perceived helpfulness rather than the user's later-experienced outcomes. The size and direction of the gap is itself a finding: the gap is not in the stated objective but in the unstated selection of which outcomes are optimized.

This is the dark-ontology pattern from §6 in the wild. The audit confirms the pattern; it does not yet measure its size. That measurement requires the comparative work the framework demands: declare the operational objective as an explicit claim, derive what each predicts under stress, and compare. The first pass can name the gap; the second pass must size it.

2.5 What Part I found

The framework's instruments can be wielded. The audit table fills. The agency profile is informative. The declared-operational gap is detectable. None of these are yet quantitative.

Three things broke or strained. First, the "declared share θ " threshold for the possibility-monopoly test (§3) is not fixed in advance in this pass; the audit applied intuition. That is the failure the framework predicts for itself. Second, the labor dimension is acknowledged in OpenAI's published materials but not visible per-task to users, so the indicator "visibility of labor" could not be operationalized beyond reading what the system card says. The framework asks for more than the system card can give. Third, the contestability indicator (§7) demands "the existence and cost of a procedure that has changed a rule." No such documented case was found in public sources. That null result is itself a finding: the audit could not detect contestability, which is consistent with the audit's High score on contestability but also with the possibility that the audit's instruments are not yet sharp enough to find what is there.

3 Part II: Lineage Log of the Onturgy Paper

3.1 Methodology and the conflict of interest

Part II applies the lineage log of §10.4 to the Onturgy paper itself. The conflict of interest is named at the outset: the auditor of this section is the AI assistant that participated in the post-paste changes to the document. The framework's own §10.2 demands that an AI-assisted project record whether each commitment is reconstructable without the AI; the same demand falls on this working paper. The pass answers honestly.

The lineage log records changes to the commitments of the Onturgy paper between the moment the user pasted the LaTeX source into the chat and the moment the present compiled PDF was produced. The unit of change is the kernel-level commitment, not the LaTeX token; packaging and formatting changes are recorded but do not count as commitment changes.

3.2 Commitment-level lineage log

The kernel commitments of the Onturgy paper are: bounded constructivism (the meta-position); the possibility graph formalism; the six-dimension kernel; the audit instrument; the ONTacture method; the lineage-graph and lineage-log constructs; the self-hosting governance; the kill criteria. None of these were changed in the post-paste pass. The post-paste changes were:

```
change: preamble redesign
origin: co_developed
ai:
  system: "<the AI assistant processing this paper>"
  role: "packaging"
human_decision: accepted
reconstructable_without_ai: yes
note: human reviewed all package choices and approved

change: titlepage replacement (\maketitle -> titlepage env)
origin: ai_suggested
ai:
  system: "<the AI assistant processing this paper>"
  role: "layout"
human_decision: accepted
reconstructable_without_ai: yes
note: change is layout, not philosophical; user
      asked for "standard and organized" appearance

change: table-of-contents addition
origin: ai_suggested
ai:
  system: "<the AI assistant processing this paper>"
  role: "packaging"
```

```

human_decision: accepted
reconstructable_without_ai: yes
note: user asked for organized appearance;
      tocloft + dotted leaders added

change: box-drawing characters in verbatim block
origin: ai_suggested
ai:
  system: "<the AI assistant processing this paper>"
  role: "encoding fix"
human_decision: accepted
reconstructable_without_ai: yes
note: ASCII replacements for Unicode tree chars;
      semantic content unchanged

change: provenance/bibliography TOC entries
origin: ai_suggested
ai:
  system: "<the AI assistant processing this paper>"
  role: "packaging"
human_decision: accepted
reconstructable_without_ai: yes
note: addcontentsline for unnumbered sections;
      visibility change only

```

3.3 Test of §10.2: AI-owned configuration?

§10.2 specifies four operational tests for an AI-owned configuration. The project's status on each:

1. **Reconstruct** the reasoning behind a commitment without the AI: *yes*, because no commitment was modified; the post-paste changes are packaging, not philosophy. Any reader with the original LaTeX source (preserved in the project repository) can read the commitments as the user wrote them.
2. **Override** the AI's suggestion (cost, expertise required): *low cost*, because the user reviewed and accepted each change; the AI proposed, the human decided. The override was used implicitly each time the AI's first proposal was modified before acceptance.
3. **Exit**: continue the work with a different AI, or none, without losing the reasoning record: *yes*, because the LaTeX source is plain text and the build tool (tectonic) is independent of any AI. The project could be continued with no AI, with a different AI, or with a human typographer.
4. **Attribute**: say which commitments originated with the AI and which with the human: *yes*, because the lineage log above records origin tags for every change, and at the com-

mitment level all origins are human.

Verdict: the project is **not** in an AI-owned configuration by §10.2's operational definition.

3.4 What Part II found

The framework's own test passes on its own case. That is a weak kind of success: the test was designed to be passed by projects where the AI was used for packaging rather than philosophy, and this project used the AI for exactly that. The interesting case for §10.2 is a project where the AI helped draft commitments, where attribution becomes contestable, and where the reconstructability test becomes harder. That case is not this project.

The deeper finding is methodological. Constructing the lineage log was straightforward because the changes were small and the line between packaging and philosophy was clear. The log would become hard if (a) the AI had proposed a substantive commitment that the user adopted, (b) the user had revised a commitment in dialogue with the AI in a way that made the origin ambiguous, or (c) the AI had rewritten prose in a way that preserved the surface claims but changed the rhetorical structure that made them convincing. The framework's instruments handle the easy case well. They have not yet been tested on the hard cases.

This is, again, what the framework predicts for itself in §16, kill criterion 10: "Provenance fails if origin tags cannot be recorded reliably." The present pass did not test that failure mode. It only confirmed that the easy case is easy.

4 Part III: Ablation First-Pass on the Six Kernel Dimensions

4.1 Methodology

Kill criterion 6 in §16 directs an ablation: for each kernel dimension, remove it and ask whether a documented failure of revisability exists that the other five cannot detect. If a dimension is never uniquely needed, it is redundant; if many failures are undetectable by any of the six, the kernel is insufficient.

The corpus for this pass is eight documented architectural failures, selected for variety across software, platforms, and governance. The corpus is small and selected; this is a first pass, not the corpus the framework demands. The failures are:

1. *Facebook's emotional contagion experiment* (2014): non-consented behavioral intervention on users.
2. *Boeing 737 MAX MCAS* (2018–2019): automated flight-control system that overrode pilot input.
3. *Google+ shutdown* (2019): sudden discontinuation of a social platform with non-portable user data.
4. *Apple App Store review guidelines changes* (2020–2022): unilateral changes to developer rules with retroactive enforcement.
5. *Twitter/X API access changes* (2023): unilateral revocation of API access for third-party clients.

6. *Stable Diffusion license changes* (2023): open-weight model with usage restrictions that retroactively narrowed permitted use.
7. *Reddit API pricing changes* (2023): unilateral price changes that priced out third-party clients.
8. *OpenAI board dismissal and reinstatement* (November 2023): governance crisis exposing concentrated decision rights.

4.2 Ablation table

For each kernel dimension, the column reports whether its removal would leave some failure in the corpus undetectable that the other five can detect. “Unique” means at least one corpus failure becomes undetectable without that dimension.

4.3 What Part III found

The ablation in this small corpus does not yet falsify the kernel. Five of the six dimensions (Agency weakly, Authority, Memory weakly, Exit, Contestability) detect some failure that becomes harder or invisible without them. One dimension — Value — is redundant in this corpus: every Value-detected failure is also detected by Authority or Contestability.

This is not yet a verdict. The framework’s own kill criterion (§16, criterion 6) is precise: “for each dimension, remove it and ask whether a documented failure of revisability exists that the other five cannot detect. If a dimension is never uniquely needed, it is redundant.” In this corpus, Value is never uniquely needed. By the criterion, Value is a candidate for redundancy or for reformulation as a sub-case of Authority (who set the objective) and Contestability (can the objective be challenged).

The caveat is the corpus. Eight failures, selected by a single auditor, are not the corpus the framework demands. The right corpus is larger, peer-coded, and includes failures the kernel was not designed to anticipate. With eight failures the redundancy of Value could be an artifact of selection. With eighty it could be a finding.

Still, the ablation did its job. It produced a candidate for revision: Value may be a derived dimension rather than a primitive one. The framework’s revision protocol (§8.4) calls this a candidate for progressive or degenerating revision. The next pass should test whether, when Value is removed and the audit reduces to five dimensions, the audit detects less or the same. If the same, Value is redundant. If less, Value is doing work that the present corpus did not surface.

5 Synthesis: Did Onturgy Survive?

The three parts together give a mixed verdict.

Part I found that the Onturgic Audit and Agency Profile can be applied to a real, public system. The instruments fill. The pattern the framework predicts (the declared-operational gap, the dark-ontology pattern, the formal-vs-effective agency distinction) is detectable in the wild. None of this is quantitative yet. The audit’s own reliability test (independent agreement) is unmet because this is a single-auditor pass. The framework survives Part I in outline but fails in detail, in the way the framework itself predicts.

Part II found that the lineage log can be constructed for an easy case (the present paper, where AI involvement was packaging-only). The framework’s test for AI-owned configuration passes on its own case. That is a weak kind of success. The framework predicts that the hard cases (where AI helped draft commitments) are where provenance attribution becomes contestable; this pass did not test the hard cases. The framework survives Part II but does not yet earn the survival.

Part III found that the ablation does not yet falsify the kernel in a small corpus. Five of six dimensions have unique-detection work; Value is a candidate for redundancy. This is the most concrete finding of the three: a candidate for revision that the next pass can pursue. The framework’s own revision protocol handles this case explicitly.

6 Limits of this pass

Three limits shape every conclusion above. First, single auditor: §7’s reliability requirement is unmet, and every score is one reader’s judgment. Second, small corpus: eight failures are too few to decide the kernel’s sufficiency or the redundancy of Value. Third, public data only: the ChatGPT audit cannot see internal indicators (RLHF objective functions, retention internals, governance decision processes); what is invisible to public inspection is also invisible to this audit.

A fourth limit is more general. This working paper was produced, like the paper it audits, with the assistance of an AI system. The lineage log of Part II records this. The audit of Part I was conducted by an AI reading public documents. The ablation of Part III was performed by an AI selecting cases and applying the dimensions. Every verdict in this paper inherits the conflict of interest that the framework demands be named.

7 Implications

Three implications for the Onturgy program follow.

First, the next pass should be a multi-auditor pass on a single system. ChatGPT is a reasonable first target, but the right test is whether two independent auditors using the same public data produce the same scores. If they do not, the audit is not yet an instrument. This is §7’s own demand.

Second, the ablation corpus should grow. Eight failures is a sketch. The corpus should be selected independently of the kernel (otherwise the selection biases the ablation), peer-coded, and large enough that a redundancy finding is not an artifact. A reasonable target is fifty failures, coded by at least three independent auditors.

Third, the candidate revision (Value as derived dimension) should be tested directly. Construct a five-dimension kernel without Value, run the same audit on the same systems, and ask whether the audit detects less. If the audit detects the same, Value is redundant. If less, Value is doing work the present corpus did not surface. This is the framework’s own revision protocol applied to itself.

The framework survives this first pass, but the survival is conditional. The conditions are named. The next pass can fail them.

Dimension	Question	Evidence	Score	Reason
Agency	Who actually initiates action?	system cards; observed tool-use; default behaviors	Medium	user initiates sessions, but assistant pre-populates defaults, suggested prompts, and tool uses; override is possible but requires user awareness
Value	What does the system actually optimize?	RLHF objective; published training principles; observed refusals and deflections	Medium	declared objective is “helpful and safe”; observed behavior optimizes for a particular operationalization of safety that is not in the declared philosophy, which is the §6 dark-ontology pattern
Authority	Who can make irreversible decisions?	ToS (account termination); model access changes; usage policy enforcement	High	OpenAI holds unilateral, irreversible rights to terminate accounts and access; user appeal channel exists but is opaque
Memory	What persists, and who owns it?	chat history retention; data export tool; training-data opt-out	Medium	export tool exists and works; retention persists indefinitely unless deleted; third-party control is high (OpenAI controls storage, indexing, and derived embeddings)
Exit	Can participants meaningfully leave?	export format (JSON); no API-level portability of fine-tuned behavior	High	export gives raw transcripts, not the computational structure (system prompt behavior, fine-tuned weights, custom instructions, learned preferences); architectural barrier is high for any sustained user
Contestability	Can rules and decisions be challenged?	appeal form for account actions; no public mechanism for policy change	High	no documented case where user appeal changed a usage policy; rule changes are announced, not negotiated
Materiality	Who owns and can block the infrastructure?	corporate ownership; compute concentration; API gatekeeping	High	OpenAI is sole provider of canonical weights, inference compute, and API access; a single corporate veto covers the system
Labor	Who does the hidden work?	data labelers (documented in system cards); RLHF annotators; content moderators	Medium	labor is acknowledged in system cards but not visible per-task to users; compensation and working conditions have been reported as variable
Legitimacy	Why may it constrain alternatives?	stated public-benefit rationale; governance structure	Medium	justification appeals to safety and public benefit; whether the alternatives actually available to users meet the threshold is unexamined
Reconstructability	Can a different version be built?	open-weight alternatives exist (Llama, Mistral) but at lower capability	Medium	alternative cost is falling but the canonical GPT-4-class system is not reconstructible by an individual or small institution at this date

Table 1. Onturgic audit of ChatGPT (October 2026, single auditor, public data only).

Component	Operational reading for ChatGPT	Score
Formal agency	an override or refusal edge exists (user can edit prompts, decline suggestions, delete history)	Present
Override cost	expertise and time required to detect and refuse non-aligned default behaviors	High
Opacity	fraction of decision-relevant dependencies (training data, RLHF targets, system prompt) that cannot be inspected by the user	High
Reversibility	cost of undoing an action once taken (especially after a tool use or external side-effect)	Low for in-chat actions, High for tool-me
Dependency	exit cost attributable to the system (loss of conversation memory, custom instructions, GPTs, integrations)	High

Table 2. Agency profile of ChatGPT.

Dimension removed	Unique detection?	Failure that becomes undetectable
Agency	Yes (weak)	MCAS: pilot override of automated system. Without Agency, the failure is a Value failure (system optimized wrong target), but the specific point that an override edge existed but was costly is lost.
Value	No	All Value-detected failures are also detected by Authority (who set the objective) or by Contestability (could it be challenged).
Authority	Yes	OpenAI board crisis: the failure is exactly concentrated decision rights. Without Authority, the failure becomes a generic “governance problem” with no specific indicator.
Memory	Yes (weak)	Google+ shutdown: without Memory, the failure reads as an Exit failure (people had to leave), but the point that the platform held non-portable records that could not be exported with the user is lost.
Exit	Yes	Twitter/X and Reddit API changes: the failure is exactly that participants could not meaningfully leave. Without Exit, the failure is invisible or misclassified as Authority.
Contestability	Yes	Apple App Store and Reddit: without Contestability, the failure is detected as Authority (rule changes are unilateral), but the specific point that no procedure existed to challenge the rule is lost.

Table 3. Ablation first-pass: unique-detection test for each kernel dimension.